

Museum Monitoring: an environmental monitoring dataspace using The Things Network, Solid, and LDES

Arthur Vercruysse¹, Ben De Meester¹, Julian Rojas¹, Dieter Suls² and Tobias Hendrickx²

¹*Ghent University - IMEC,*

²*Fashion Museum Antwerp, Nationalestraat 28, 2000 Antwerpen*

Abstract

Museums increasingly deploy environmental monitoring systems (e.g., data loggers and sensor networks) to support preservation of cultural heritage objects. However, monitoring data is often captured and accessed through proprietary vendor software and local infrastructures, which makes reuse, integration, and controlled cross-institution sharing difficult—particularly in object loan scenarios. This paper reports on MuMo (Museum Monitoring), a three-year applied research project that explored how dataspace principles can be applied to environmental monitoring in real museum settings. Rather than replacing existing systems, MuMo extends a legacy monitoring dashboard with semantic data modeling, Linked Data Event Streams, and Solid-compliant data access management. We present the system design: environmental monitoring data is semantically described and published as access controlled streams of fragments, and the semantic links between these streams allow clients to prune irrelevant branches while providers can restrict access at natural boundaries (e.g., per group or sensor). Further, we describe how MuMo is used in practice through a set of in-use scenarios, where the system enables (1) decentralized data publication for longitudinal analysis of environmental conditions, (2) selective cross-institutional data sharing during object loans, and (3) client-side data integration and aggregation across independent deployments. In practice, this dataspace-oriented deployment reveals design trade-offs while preserving institutional autonomy: museums control their data and authorization policies end-to-end (sufficiently with group-based access control), supporting trusted data sharing across organizational boundaries. The deployed system and its insights are thus relevant to a broad range of (Digital Humanities) projects that involve long-running data integration under distributed governance.

Keywords

Digital Humanities, Dataspaces, Solid, LDES, Museum Monitoring

1. Introduction & Motivation

In cultural heritage, environmental parameters such as temperature, humidity, and light exposure directly affect the preservation of artworks [12]. Thus, museums increasingly require digital infrastructures for long-term environmental monitoring and traceable reporting [Laborda et al. [12]][15].

A wide range of sensing technologies for environmental monitoring already exists and is actively used in museum practice, including data loggers and long-running wireless sensor network (WSN) deployments in museum buildings [16]. However, these systems typically operate as stand-alone solutions, tightly coupled to proprietary (and legacy) dashboards or local infrastructures [2]. As a consequence, monitoring data remains siloed: it is difficult to combine measurements across installations [1], to align them with other data sources, or to selectively share specific sets of measurements with external partners through access control mechanisms. This fragmentation limits the analytical value of the data and complicates reporting and thus collaboration between institutions.

Woodstock'22: Symposium on the irreproducible science, June 07–11, 2022, Woodstock, NY

*Corresponding author.

†These authors contributed equally.



© 2022 Copyright for this paper by its authors. Use permitted under Creative Commons License Attribution 4.0 International (CC BY 4.0).

The limitations of siloed monitoring become particularly visible in inter-museum loan scenarios: a lending museum lends a particular cultural heritage object for a specific period of time to another borrowing museum. In this scenario, borrowing museums are typically contractually obliged to give the lending institution access to the borrowing museum’s environmental data, but only for a specific subset of sensors and a limited time period, relevant to the loan. As a result, monitoring data must not only be collected reliably, but also interpreted, analyzed, and—under specific conditions—shared across institutional boundaries. Existing practices often rely on ad-hoc data exports or manual reporting, leading to delays, duplication, and reduced transparency in conservation oversight.

Recent work has similarly highlighted how the increasing digitization of museum collections exposes structural limitations in legacy monitoring and documentation infrastructures, particularly with respect to long-term traceability, institutional accountability, and resilience against security incidents [17].

This paper reports on MuMo (Museum Monitoring), a three-year applied research project that explored how dataspace principles can be applied to environmental monitoring in real museum settings, namely prioritizing incremental integration, heterogeneous sources, and governance. Rather than proposing new sensing technologies or replacing existing platforms, MuMo focused on breaking down data silos by integrating semantic technologies, decentralized data publication, and access control mechanisms with existing systems and workflows.

This paper makes three contributions: (1) an in-use account of deploying dataspace principles for environmental monitoring in museums, (2) a concrete architectural integration of semantic modeling, event-based publication, and decentralized governance with legacy systems, and (3) empirically grounded lessons on aligning access control and data sharing mechanisms with institutional practice.

After describing the X in Section Y, we [...]

2. Context of Use

The Museum Monitoring (MuMo) project was a three-year cultural heritage project¹, funded by the Flemish Government from 2022 to 2025. Its applied research objective was to develop both open-hardware monitoring devices and a dedicated data space to support museums in the long-term tracking of environmental conditions surrounding collection objects.

2.1. Operational Setting

MuMo has been deployed in real museum environments and the sensors were placed near artworks or storage locations and continuously measured environmental parameters. Due to the physical and organizational constraints of museum spaces, the system had to function with minimal intervention once deployed. To accommodate this requirement, custom ultra-low-power sensing hardware was developed, designed to operate for approximately one year without recharging.

Measurements captured by the MuMo monitoring devices were transmitted via a LoRaWAN signal to off-the-shelf gateways and routed through The Things Network, where they were then ingested into an existing (legacy) monitoring dashboard². This dashboard became the primary operational interface for museum staff and therefore strongly shaped how data could be accessed, interpreted, and shared.

2.2. Legacy Dashboard Constraints

The existing dashboard provided:

¹<https://www.vlaanderen.be/cjm/nl/cultuur/cultureel-erfgoed/subsidies-cultureel-erfgoed/projects/subsidies-cultureel-erfgoed/internationale-landelijke-en-bovenlokale-cultureel-erfgoedprojecten>

²<https://www.thethingsnetwork.org/docs/lorawan/architecture/>

- User and group management, including recursively defined groups
- Node management corresponding to deployed sensors
- Basic visualization through simple time-series graphs
- Configure alerts and notification when measured values cross a certain threshold.
- Export data ranges to CSV

Access control was **group-based and coarse-grained**, sufficient for internal monitoring within a single institution but not designed for cross-institutional collaboration. Importantly, the dashboard could not be replaced or fundamentally re-engineered within the scope of the project, requiring MuMo to operate within these constraints.

2.3. Cross-Institutional Access and Loans

A central real-world use case was **the monitoring of artworks during loans between museums**. Lending institutions often require insight into the environmental conditions under which an object is kept while on loan, without being granted broader access to the borrowing museum’s infrastructure or internal data.

In practice, this requirement translates into:

- access limited to a subset of sensors,
- bounded by the duration of the loan,
- manageable by non-technical museum staff.

Although often described as “fine-grained access control,” the actual operational need is **group-level authorization**: sensors associated with a specific loan can be grouped, and access to that group can be granted temporarily to another institution.

2.4. Requirements for Data Sharing

These scenarios impose several practical requirements:

- Data must be shareable across **institutional boundaries**.
- Access must respect existing group-based authorization practices.
- Users must be able to access data originating from multiple museum setups.
- Identity management cannot assume a single central authority.

These requirements motivated an architecture in which monitoring data could be published, discovered, and accessed across organizational boundaries, while remaining compatible with legacy systems and established museum workflows.

3. System and Approach Overview

This section presents the architecture of MuMo and explains the main design choices behind its approach to cross-institutional environmental monitoring. Rather than replacing existing museum monitoring infrastructures, MuMo adds a semantic publication and governance layer that enables decentralized data access and controlled sharing across institutions. We first outline the overall system composition and data flow, and then describe the key building blocks—semantic modeling, Linked Data Event Streams for continuous publication, and Solid-based governance—showing how they work together to support both longitudinal analysis and collaboration scenarios such as object loans.

3.1. Dataspace Architecture Based on Solid

MuMo adopts a **dataspace-oriented architecture** based on Solid: a decentralization framework in which data storage [+++ solid protocol], identity [+++ solid-oidc], and access control [5] are decoupled from individual applications, enabling data to be reused by multiple clients and actors without funneling all data through a centralized service [4, 14, 23]. This approach aligns with museum practice: institutions are unwilling to relinquish control over their monitoring infrastructure, yet must be able to selectively share data during collaborations such as object loans.

Some recent research emphasizes cryptographic immutability and tamper resistance, by coupling sensor infrastructures with hardware-based attestation and private distributed ledgers [17]. Within MuMo, we prioritized institutional autonomy and operational feasibility by enabling decentralized data publication and access control aligned with existing museum workflows.

In MuMo, Solid Pods act as **institutional data endpoints** rather than user-centric storage, enabling long-lived publication of monitoring data that remains independently governed.

3.2. Continuous Data Publication with Linked Data Event Streams

Environmental monitoring produces sensor data: **continuous, append-only data** that grows over time and is rarely modified retroactively. To reflect this, MuMo publishes monitoring data using **Linked Data Event Streams (LDES)** [7, 22].

LDES is a pattern that is increasingly adopted for interoperable publication of evolving datasets, enabling replication and synchronization across organizational boundaries without centralized query services [7].

An LDES is typically published as a tree-based hierarchy of fragments: each fragment links to related fragments through semantic relations. These relations allow consumers to navigate to relevant subtrees without requiring them to scan the full stream. This publication pattern enables consumers to (1) retrieve historical data incrementally, (2) stay synchronized with newly produced observations, and (3) avoid repeated querying of centralized services.

Choosing how events are grouped into such a fragment tree is known as the fragmentation strategy. In MuMo, observations are fragmented along the operational dimensions that curators and technicians naturally work with: group, sensor, and time. Concretely, we publish observations under a hierarchical path of the form group-x/sensor-x/year/month/day. This strategy keeps fragments small and stable over time, while preserving a semantic navigation structure (e.g., only follow fragments for sensor X within the selected time window).

This publication model proved particularly suitable in a cross-institutional setting, as it allows data consumers to process only the subsets of data they are authorized to access, without requiring the data provider to offer tailored query endpoints [22].

Prior work has demonstrated the feasibility of storing and consuming LDES event sources within the Solid ecosystem, supporting the compatibility of event-stream publication with Solid-based access control [18].

3.3. Semantic Representation of Sensors and Observations

All data managed by the legacy dashboard is transformed into a semantic representation, including both **environmental observations** (e.g., temperature, humidity, and light exposure) and **sensor configuration metadata** (i.e., group membership, location, and [+++]). This separation reflects two different kinds of change:

- observations evolve continuously over time,
- sensor configurations evolve discretely when sensors are moved or reconfigured.

The semantic representation builds on established models for describing sensors and observations, most notably the W3C Semantic Sensor Network ontology (SSN/SOSA) [6, 9], which provides a shared conceptual framework for sensors, observations, and their relationships.

SSN/SOSA was selected in particular for its explicit separation between sensors, observations, and the features of interest being observed, enabling sensor redeployment and contextual change to be represented without altering historical observations [6].

In MuMo, these concepts are instantiated using OSLO vocabularies [3], a Flemish Linked Data standardization framework that profiles and reuses international semantic standards for cross-organizational data exchange.

Sensor configuration metadata is published as a **versioned event stream**, this allows consumers to reconstruct the context in which observations were produced. Group membership and location are encoded explicitly in the sensor configuration metadata, which enables downstream systems to independently interpret authorization rules correctly.

3.4. Group-Based Access Control as a Practical Design Choice

Access control in MuMo deliberately mirrors the **group-based authorization model** already in use in the legacy dashboard. Rather than introducing fine-grained authorization at the level of individual observations, access is granted at the level of sensor groups.

At the Solid protocol level, this aligns with authorization mechanisms such as Web Access Control (WAC), which allow servers to enforce access rules on resources and their associated representations [4, 5].

This choice was guided by museum practice: for loan scenarios, institutions require access to all monitoring data related to a specific object or location over a defined period. Prior work on access control has repeatedly highlighted that fine-grained authorization models, while expressive, introduce significant configuration and usability complexity in operational settings, particularly for non-technical users [10]. Group-level access was therefore found to be both sufficient and manageable, avoiding complexity that would hinder adoption.

Authorization policies defined in the legacy system are reflected in access constraints on published data, allowing cross-institutional sharing without introducing a new access management interface.

The fragmentation strategy used for the LDES publication directly supports this governance model. Since fragments are organized as group-x/sensor-x/year/month/day, Solid’s access control mechanisms (e.g., WAC) can be applied at different levels of the fragment tree. In the deployed system, authentication impacts the first level of the hierarchy: users gain access to the subtree corresponding to a group, which implicitly grants access to all sensors and time slices contained within that group.

At the same time, this structure makes it explicit that finer-grained authorization is technically possible if required by future workflows: access could be restricted at sensor level (by granting only the corresponding sensor subtree) or at temporal level (by granting only fragments for a selected time range). MuMo intentionally does not activate these options to remain aligned with what museum staff can practically configure in the authoritative legacy dashboard.

Although the underlying fragmentation strategy and event-based publication model would technically allow finer-grained access control (e.g., at the level of individual days or measurement types), MuMo intentionally limits authorization to group-level permissions in order to remain aligned with the configuration mechanisms and functional requirements of the legacy dashboard.

3.5. Decentralized Consumption and Aggregation

Because each MuMo deployment publishes its data independently, consumers may need to combine data from multiple sources. MuMo supports this through a client-side dashboard that




Figure 1: Figure 1: deployment overview

retrieves sensor descriptions first, determines the user’s authorized scope, and then incrementally consumes the relevant observation streams.

This approach avoids centralized aggregation while still enabling a unified user experience, reinforcing the dataspace principle that **integration occurs at the point of use**.

3.6. Deployment Components and Access-Control Flow

Figure 1 shows the four runtime components that together bridge the legacy monitoring dashboard with Solid-based identity and authorization. MuMo deliberately keeps the legacy dashboard as the authoritative interface for non-technical staff, while reflecting its group-based authorization decisions into Solid’s Web Access Control (WAC) layer. This is consistent with the project’s broader constraint that the dashboard cannot be replaced and already provides the operational concepts of users, groups, and recursively defined group hierarchies.

3.6.1. Component 1 — PHP dashboard (admin & group management)

The PHP dashboard is the operational control plane. Administrators manage:

- Groups and group hierarchies (including recursive nesting),
- User-to-group assignments, and
- Cross-institution access by granting group access to WebIDs.

A WebID may originate from the deployment’s *local Solid Identity Provider (IdP)* (Component 4) or from an external provider such as Inrupt. From the dashboard’s perspective, a WebID is simply the stable identifier used to grant or revoke access at group level—matching how museum staff already reason about “who gets access to which loan group”.

3.6.2. Component 2 — Solid Pod hosting LDES resources (data plane)

The Solid Pod is the enforcement point for data access. It hosts the published resources (including the LDES fragment tree) and applies WAC authorization before serving any protected resource. Conceptually, this is where an incoming request is interpreted in terms of which protected subtree (group/sensor/time path) is being accessed, and whether the requester’s authenticated WebID is allowed to access it.

3.6.3. Component 3 — ACL generator (policy materialization service)

The ACL generator is the bridge between the dashboard’s group model and Solid’s resource-level authorization. It receives (synchronized) configuration derived from Component 1:

- For each configured *WebID*: the *groups* it belongs to;
- For each group: its *position in the group hierarchy* (to support inherited/recursive membership).

From this input, the ACL generator *materializes WAC ACL documents on demand*. Operationally, its central function is:

given a request targeting a particular group (e.g., “group 4”), return the effective ACL representation for that group.

Rather than requiring administrators to maintain ACL files manually (which would be error-prone and mismatched with existing workflows), the generator ensures the Solid-side authorization view remains a faithful projection of the dashboard configuration.

3.6.4. Component 4 — Solid Identity Provider (IdP)

The Solid IdP provides authentication for WebIDs under the local deployment’s authority. However, MuMo does *not* require a single centralized identity authority: when a user presents a WebID from a different IdP, authentication can be performed against that remote IdP as well (subject to standard Solid/OIDC behavior). This keeps identity decentralized while still allowing the dashboard to manage access using WebIDs as portable identifiers.

3.7. End-to-end request flow

1. **Admin configuration (control path):** An administrator updates group membership in the PHP dashboard by associating WebIDs with groups and maintaining the relevant group hierarchy.
2. **Sync to policy service:** These settings are synchronized to the ACL generator as (WebID → groups) plus the group hierarchy.
3. **User data request (data path):** A client requests an LDES resource from the Solid Pod.
4. **ACL resolution:** Before serving the resource, the Solid Pod determines which ACL scope applies (effectively: which group subtree is being accessed?) and consults the ACL generator to obtain the effective ACL document for that scope.
5. **Authentication + authorization:** The Solid Pod validates the requester’s identity via the appropriate IdP (local or remote) and evaluates whether the authenticated WebID is granted access by the generated ACL.
6. **Response:** If authorized, the Solid Pod serves the requested LDES resource; otherwise it denies access.

4. In-Use Scenarios

This section illustrates how the MuMo dataspace architecture is used in practice by museum professionals to analyze, share, and contextualize environmental monitoring data.

4.1. Scenario 1: Analyzing Environmental Conditions Over Time

The primary use of the advanced dashboard is to support museum staff in assessing the long-term “health” of collection objects by analyzing environmental conditions over time. Users interact with the system through a web-based dashboard that allows them to filter and combine data based on:

- location (group),
- sensor or node,
- type of measurement (e.g., temperature, humidity),
- time constraints.

Using these filters, users can construct queries that follow an artwork throughout its lifecycle. For example, a curator can analyze how environmental conditions evolved while an object was stored in one room, then exhibited in another, and later placed in temporary storage. Because sensor configuration changes are published as versioned semantic descriptions, the system can correctly associate observations with their deployment context at each point in time.

A key practical benefit of the Linked Data Event Streams (LDES) publication model is that **data filtering occurs before retrieval**. Rather than fetching all historical measurements and filtering client-side, the dashboard retrieves only the relevant fragments based on semantic relations and temporal constraints. For instance, when analyzing conditions in a specific year, only the corresponding fragments are accessed, avoiding unnecessary data transfer and improving responsiveness.

This is enabled by the fact that the event stream is published as a semantically linked fragment tree. For a given query, the dashboard can follow only those fragment relations that match the selected group, sensor, and time window, and prune all other subtrees. This makes selective data access a navigation problem rather than a centralized query problem, which fits cross-institutional settings where providers should not have to expose custom query endpoints.

This incremental and selective access is essential for long-running monitoring installations, where datasets grow continuously and may span multiple years.

4.2. Scenario 2: Cross-Institutional Access During Loans

A second, critical scenario concerns the monitoring of artworks during loans between museums. Lending institutions typically require access to environmental data from the borrowing museum to ensure that conservation conditions meet agreed standards, while borrowing institutions must retain control over their broader monitoring infrastructure.

In MuMo, this scenario is supported through **group-based access control aligned with Solid identities**. For a loan, the borrowing museum creates a dedicated group in the legacy dashboard and associates the relevant sensors with that group. Access to this group is then granted to specific Web-based identifiers belonging to the lending institution.

Because access control is enforced at the level of published data fragments, external users can authenticate using their own WebIDs and access only the data streams corresponding to the loan group. No centralized user management or data replication is required. Once the loan period ends, sensors should be taken out of the group, so new data is not shared.

This approach enables temporary, fine-grained sharing of monitoring data across institutional boundaries while remaining manageable for museum staff and compatible with existing workflows.

4.3. Scenario 3: Integrating Multiple Data Sources

In addition to supporting analysis within a single monitoring deployment, the advanced dashboard demonstrates the ability to **combine data from multiple independent MuMo data sources**.

Each MuMo deployment publishes its monitoring data and sensor descriptions independently, yet follows the same semantic representation and event-based publication model.

In practice, this allows users to access and analyze data originating from different museum setups within a single interface. For example, a user may compare environmental conditions across multiple exhibition spaces or institutions, provided they have the appropriate access rights. Because sensor metadata and observations are published as Linked Data Event Streams, the dashboard can discover available sensors, determine authorization, and incrementally retrieve data from multiple sources without requiring centralized aggregation.

Beyond this demonstrated functionality, the same mechanisms also enable the **conceptual integration of external data sources** that are not part of the MuMo project, such as weather station measurements. Since both sensor metadata and observations are modeled using shared semantic standards, incorporating additional event streams would not require changes to the underlying architecture. While such external integrations have not yet been deployed, they directly informed the design of the system and illustrate how the dataspace approach supports extensibility and reuse.

4.4. Generalizing the scenarios beyond monitoring systems

While the scenarios above focus on environmental monitoring data, they also reveal a more general pattern regarding the role of contextual information in enabling downstream analysis. They rely on contextual information such as temporal validity, object location, and sensor deployment. In MuMo, this information was obtained from the monitoring infrastructure and its associated configuration data. However, the underlying mechanism illustrated by these scenarios does not depend on the specific system from which such context originates. In practice, museums already maintain similar forms of structured contextual information in a variety of collection management and repository systems, independent of environmental monitoring.

Contemporary collection management systems and digital repositories often manage object metadata, temporal associations, and organizational context, and may support structured export mechanisms. Examples include repository platforms such as Omeka-S or digital repository services such as D-RaaS [20].

When such contextual information is made available in a structured and time-aware form, it becomes possible to derive queries that are not explicitly supported by the originating CMS. Instead, these queries emerge from the combination of temporal relations, location changes, and object associations at the point of use.

MuMo demonstrates this principle in the domain of environmental monitoring, but the same mechanism can be applied more broadly to contextual data maintained by other institutional systems. This suggests that dataspace-oriented approaches can unlock additional analytical value from existing CMS infrastructures without requiring their replacement or redesign.

5. Observations and Lessons Learned

The in-use scenarios presented above highlight how specific architectural and modeling decisions shaped the practical use of MuMo in museum monitoring contexts. Rather than evaluating individual technologies in isolation, this section reflects on how these choices interacted with real-world constraints and practices.

5.1. Solid as an Institutional Boundary Mechanism

Across all scenarios, Solid played a central role in enabling data sharing without centralization. By treating Solid Pods as **institution-level data endpoints** rather than user-centric storage, MuMo aligned naturally with museum governance structures. Institutions retain control over their monitoring data while still enabling external access when required.

This was particularly evident in the loan scenario (Scenario 2), where Solid identities (WebIDs) allowed users from different institutions to authenticate without relying on a shared user database. Access could be granted and revoked by modifying group membership, supporting temporary collaborations without introducing new identity management workflows.

Lesson learned: Solid is most effective in this context when used to represent **organizational boundaries and responsibilities**, rather than as a personal data store.

5.2. Linked Data Event Streams for Scalable Analysis

The use of Linked Data Event Streams proved essential for handling the continuous and growing nature of monitoring data. In the analysis scenario (Scenario 1), LDES enabled the dashboard to retrieve only the relevant subsets of data based on semantic relations and temporal constraints, avoiding the need to fetch complete datasets and filter client-side.

Lesson learned: The combination of LDES and Solid is particularly effective because fragmentation simultaneously serves performance and governance needs. Fragment trees enable clients to prune irrelevant subtrees for a given query, while Solid access control can be enforced on the same subtree boundaries. This allows selective cross-institutional access without requiring providers to implement bespoke query endpoints or replicate data into centralized stores.

This incremental access model also facilitated multi-source analysis (Scenario 3), where data from multiple independent MuMo deployments could be consumed and combined without requiring centralized aggregation or bespoke query endpoints.

Lesson learned: Event-based publication is particularly well suited for long-running Digital Humanities data collection, where datasets evolve continuously and selective access is more important than expressive querying.

5.3. Group-Based Access Control as a Deliberate Design Choice

MuMo deliberately adopts **group-level access control**, even though finer-grained authorization is technically feasible within the chosen architecture. The use of Linked Data Event Streams combined with a hierarchical fragmentation strategy makes it possible to enforce access constraints at different levels of granularity, such as individual days or, with minor adaptations to the publication pipeline, per measurement type.

In practice, these options were not activated. The legacy dashboard that remains authoritative for configuration and access management does not provide mechanisms to express such fine-grained permissions. As a result, only those authorization concepts that are meaningful and manageable within existing museum workflows were reflected in the generated access control configurations.

This decision was particularly evident in the loan scenario, where group-level access proved sufficient to grant lending institutions visibility into relevant monitoring data without exposing unrelated information. Introducing finer-grained access control would have increased technical complexity without corresponding benefits for end users.

Lesson learned: In applied Digital Humanities settings, the appropriate granularity of access control is determined less by technical capability than by the expressive power of existing organizational tools and practices. Designing for feasible access control is often more valuable than designing for maximal access control.

5.4. Semantic Modeling as an Enabler of Integration

Semantic representations based on shared observation models enabled MuMo to separate **data production** from **data consumption**. Sensor descriptions, deployment context, and observations could be published independently, versioned over time, and interpreted consistently across systems.

This separation was critical in Scenario 1, where sensor relocations needed to be reflected correctly in longitudinal analyses, and in Scenario 3, where data from multiple sources could be combined as long as they adhered to the same semantic structures.

Lesson learned: The value of semantic modeling in practice lies less in expressiveness and more in its ability to support evolution, reuse, and alignment across independently managed systems.

5.5. Dataspace Integration at the Point of Use

Finally, the combination of Solid, LDES, semantic modeling, and client-side aggregation reflects a broader dataspace principle: **integration happens at the point of use**, not through centralized infrastructure.

Similar architectural patterns can be observed in other domain-specific data space initiatives, such as the Flanders Smart Data Space, where semantic standardization and Linked Data Event Streams are combined to support decentralized data publication across organizational boundaries [13].

More broadly, European data space initiatives emphasize federated trust, identity, and governance mechanisms to enable cross-organizational data exchange while preserving sovereignty, aligning with the boundary-oriented role Solid plays in MuMo [11, 19].

Rather than enforcing a single global schema or repository, MuMo enables institutions to publish data independently and allows consumers to integrate only what they need, when they need it. This approach proved compatible with legacy systems and institutional autonomy, both of which are common in Digital Humanities settings.

Lesson learned: Dataspace architectures are particularly well suited to DH environments where data is distributed, governance is decentralized, and collaboration is episodic.

6. Impact for the Digital Humanities Community

Many Digital Humanities projects rely on data that is collected continuously over long periods of time and embedded in local infrastructures, such as sensor data, annotations, or evolving metadata. In practice, such data often remains siloed within project-specific systems, limiting its reuse, comparability, and shareability across institutional boundaries.

Digital Humanities researchers have increasingly emphasized that the sustainability of digital resources is not only a technical issue but also an institutional and socio-technical one, shaped by funding horizons, maintenance practices, and governance constraints [21, 8].

The experience reported in this paper suggests that **dataspace-oriented approaches** are well suited to such settings. By allowing data to remain under the control of the institution that produces it, while enabling integration at the point of use, dataspace provide a viable alternative to centralized platforms. This is particularly relevant for DH collaborations involving multiple partners with heterogeneous governance structures and long-lived legacy systems.

The use of **semantic representations as infrastructural connectors**, rather than as expressive modeling exercises, further supports interoperability in constrained environments. Lightweight semantic alignment enables independently managed datasets to be combined, extended, and reinterpreted over time, without requiring uniform tooling or deep semantic expertise from end users.

Access control emerged as a critical dimension of data reuse. Rather than pursuing maximal technical expressiveness, the project highlights the importance of aligning authorization mechanisms with existing organizational practices. Group-based access control, when grounded in curatorial workflows, can provide sufficient protection and flexibility while remaining understandable and manageable in daily operation.

Finally, breaking down data silos enables new forms of longitudinal and comparative analysis that are difficult to achieve with isolated systems. The ability to combine data from multi-

ple independent sources—under controlled access conditions—opens opportunities for richer contextualization and collaboration in a wide range of Digital Humanities scenarios.

Together, these observations point toward a pragmatic path for DH infrastructures that emphasizes decentralization, semantic interoperability, and governance-aware design over comprehensive but fragile integration solutions.

7. Conclusion and Future Work

This paper reported on MuMo (Museum Monitoring), a three-year applied research project that explored the use of dataspace principles for environmental monitoring in museum practice. By integrating semantic data publication, event-based access, and decentralized governance with existing monitoring infrastructure, MuMo addressed practical challenges related to long-running data collection, cross-institutional collaboration, and controlled data sharing.

A central outcome of the project is the demonstration that dataspace-oriented architectures can be deployed incrementally alongside legacy systems. Rather than requiring centralized platforms or uniform tooling, MuMo shows how interoperability and data reuse can be achieved while preserving institutional autonomy and established workflows. The project further illustrates that design choices grounded in curatorial practice—such as group-based access control—can be both sufficient in practice and easier to sustain than more fine-grained alternatives.

Several limitations remain. Access control is currently restricted to group-level permissions, reflecting the expressive capabilities of the legacy dashboard rather than technical constraints of the dataspace architecture. Future work could explore finer-grained authorization mechanisms, such as temporal or measurement-level access, should corresponding user requirements arise. In addition, while the system supports aggregation across multiple MuMo deployments, integrating external data sources beyond the project has not yet been explored in deployed settings.

In conclusion, MuMo provides an in-use perspective on how Solid-based data governance, Linked Data Event Streams, and lightweight semantic modeling can support practical Digital Humanities workflows. The project highlights the value of aligning technical architectures with organizational realities and contributes concrete lessons for the design of interoperable, sustainable DH data infrastructures.

References

- [1] Loan-related data entry survey – final results (13 january 2024), January 2024. URL <https://cidoc.mini.icom.museum/working-groups/documentation-standards/loan-related-data-entry-survey-2024-01-13/>.
- [2] Environmental monitoring, December 2025. URL https://www.conservation-wiki.com/wiki/Environmental_Monitoring. MediaWiki page. Last edited 30 December 2025.
- [3] Raf Buyle, Laurens De Vocht, Mathias Van Compernelle, Dieter De Paepe, Ruben Verborgh, Ziggy Vanlischout, Björn De Vidts, Peter Mechant, and Erik Mannens. Oslo: Open standards for linked organizations. In *Proceedings of the international conference on electronic governance and open society: Challenges in Eurasia*, pages 126–134, 2016.
- [4] Sarven Capadisli, Tim Berners-Lee, Ruben Verborgh, and Kjetil Kjernsmo. Solid protocol, December 2022. URL <https://solidproject.org/TR/2022/protocol-20221231>.
- [5] Story Henry Capadisli Sarven and Tim Berners-Lee. eb access control (wac). Technical report, MIT CSAIL, 2024. URL <https://solidproject.org/TR/wac>.
- [6] Michael Compton, Payam Barnaghi, Luis Bermudez, Raul Garcia-Castro, Oscar Corcho, Simon Cox, John Graybeal, Manfred Hauswirth, Cory Henson, Arthur Herzog, et al. The ssn ontology of the w3c semantic sensor network incubator group. *Journal of Web Semantics*, 17:25–32, 2012.

- [7] European Commission. Linked data event streams (ldes) | semic support centre | interoperable europe portal, 2026. URL <https://interoperable-europe.ec.europa.eu/collection/semic-support-centre/linked-data-event-streams-ldes>. Accessed: 2026-01-20.
- [8] Katrina Simone Fenlon. Sustaining digital humanities collections: challenges and community-centred strategies. *International Journal of Digital Curation*, 15(1):13–13, 2020.
- [9] Armin Haller, Krzysztof Janowicz, Simon J. D. Cox, Maxime Lefrançois, Kerry Taylor, Danh Le Phuoc, Joshua Lieberman, Raúl García-Castro, and Claus Stadler. Semantic sensor network ontology. W3c recommendation, W3C, 2017. URL <https://www.w3.org/TR/vocab-ssn/>.
- [10] Vincent C Hu, David Ferraiolo, Rick Kuhn, Adam Schnitzer, Kenneth Sandlin, Robert Miller, Karen Scarfone, et al. Guide to attribute based access control (abac) definition and considerations. *NIST special publication*, 800(162):1–54, 2014.
- [11] International Data Spaces Association. Ids reference architecture model - IDS-RAM 4.0. Technical report, International Data Spaces Association, feb 2023. URL <https://docs.internationaldataspaces.org>. Accessed: 2026-01-20.
- [12] Jaime Laborda, Ana María García-Castillo, Ricardo Mercado, Andrea Peiro-Vitoria, and Angel Perles. From concept to validation of a wireless environmental sensor for the integral application of preventive conservation methodologies in low-budget museums. *Heritage Science*, 10(1):1–17, 2022.
- [13] Dwight Van Lancker, Steven Logghe, Julián Andrés Rojas, Annelies De Craene, Ziggy Vanlshout, and Pieter Colpaert. Semantic and technically interoperable data exchange in the flanders smart data space. In Gianluca Demartini, Katja Hose, Maribel Acosta, Matteo Palmonari, Gong Cheng, Hala Skaf-Molli, Nicolas Ferranti, Daniel Hernández, and Aidan Hogan, editors, *The Semantic Web – ISWC 2024*, pages 289–303, Cham, 2025. Springer Nature Switzerland. ISBN 978-3-031-77847-6.
- [14] Essam Mansour, Andrei Vlad Sambra, Sandro Hawke, Maged Zereba, Sarven Capadisli, Abdurrahman Ghanem, Ashraf Aboulmaga, and Tim Berners-Lee. Solid: A platform for decentralized social applications based on linked data. Technical report, MIT CSAIL, 2016. URL <https://solidproject.org/TR/solid>.
- [15] Stefan Michalski. The ideal climate, risk management, the ashrae chapter, proofed fluctuations, and towards a full risk analysis model. *Experts roundtable on sustainable climate management strategies*, pages 1–19, 2007.
- [16] Laura M Rodríguez Peralta and Lina M Brito. An integrating platform for environmental monitoring in museums based on wireless sensor networks. *International Journal on Advances in Networks and Services*, 3(1/2):114–124, 2010.
- [17] Dian Marie Ross. *DIGITAL TWINS, HARDWARE ORACLES AND DISTRIBUTED LEDGER TECHNOLOGIES*. PhD thesis, UNIVERSITY OF BRITISH COLUMBIA (Vancouver), 2024.
- [18] Wout Slabbinck, Ruben Dedecker, Sindhu Vasireddy, Ruben Verborgh, and Pieter Colpaert. Linked data event streams in solid ldp containers. In *Managing the Evolution and Preservation of the Data Web 2022*, volume 3339, pages 28–35. CEUR, 2023.
- [19] Hubert Tardieu. *Role of Gaia-X in the European Data Space Ecosystem*, pages 41–59. 07 2022. ISBN 978-3-030-93974-8. doi: 10.1007/978-3-030-93975-5_4.
- [20] Lefteris Tsipi, Demosthenes Vouyioukas, Georgios Loumos, Antonios Kargas, and Dimitrios Varoutas. Digital repository as a service (d-raas): Enhancing access and preservation of cultural heritage artifacts. *Heritage*, 6(10):6881–6900, 2023.
- [21] Joanna Tucker. Facing the challenge of digital sustainability as humanities researchers. *Journal of the British Academy*, 10:93–120, 2022.
- [22] Dwight Van Lancker, Pieter Colpaert, Harm Delva, Brecht Van de Vyvere, Julián Rojas Meléndez, Ruben Dedecker, Philippe Michiels, Raf Buyle, Annelies De Craene, and Ruben Verborgh. Publishing base registries as linked data event streams. In *International Conference on Web Engineering*, pages 28–36. Springer, 2021.

- [23] Verstraete, Melanie and Verbrugge, Sofie and Colle, Didier. Solid : enabler of decentralized, digital platforms ecosystems. In *ITS 31st European Conference 2022, Proceedings*, page 19, 2022. URL <https://itseurope.org/2022/>.